

Homework: Benchmarking Flops and Iops

Erik Saule

In this homework you will learn to:

- Interpret the technical specification of a processor
- Predict the flops and iops performance of a processor based on specification
- Write code that maximize the utilization of the processor
- Measure the flops/iops performance of a code
- Use the Model-Code-Measure-Reflect performance approach

The goal of this exercise is to try to reach the maximum rate of floating point operation a processor can perform.

1 Model

Question: What is a compute node of Centaurus composed of? What kind of processor it has? Which architecture is the processor built on? (For instance, some core i7 are based on the ivy-bridge architecture, some on the haswell architecture, ...)

Question: What is the maximum number of floating point operations this machine can perform per second?

Question: What is the maximum number of integer operations this machine can perform per second?

Hint: Remember what matters: frequency, core count, but also instruction set, the number of execution units within the cores, and what they can do at once.

2 Benchmark

Let's write two programs one to reach peak Flops and one to reach peak Iops.

Note, that the purpose of this exercise is to get the highest amount of operations performed. There is no restriction on what operations are performed or even if it makes sense to perform them. The problem is really "do as many flops/iops as you can".

Question: Write a code to get peak Flops.

Question: Write a code to get peak Iops.

Hint: Often when writing this kind of microbenchmark, the computational kernel is not doing anything useful and the compiler might see that and optimize it out. Putting the computation in its own function in a different file often prevents the compiler from realizing that.

Measure the performance that you obtain and see if it matches your expectation.

Question: How many Flops did the code achieve?

Question: How many Iops did the code achieve?

Question: Does that match expectation? Where does the discrepancy come from?

Question: Can you do better? (The closer the benchmark is to the model the more points.)

A Figuring out expected performance

Check out the reference slide of the lecture for external pointers.

B Measuring time

Remember when measuring time to measure a long period to remove startup overheads and innacuracy of the measure. In C, `gettimeofday` is typically the time function you want to use. In C++, use the new chrono features <http://en.cppreference.com/w/cpp/chrono>.

Chrono is great! Use Chrono!

C Using multiple threads with OpenMP

You might need to use multiple threads. The simplest way of doing that is using OpenMP, which is a C/C++ extension supported by most compilers. You can find a simple openmp code that starts as many threads as execution contexts in the slides. Remember to compile with `-fopenmp`.

D Manual vector instruction

You might also need to use compiler built-in/intrinsic functions to manually leverage SIMD operations. In most compilers they become available when you include `immintrin.h` but you might need to pass special flags to enable them (such as `-march=native` or `-mavx`).

Check out reference slide of lecture for links documenting these functions.

AVX2 Instructions for Visual Studio. To enable AVX2, go to Project Properties – Configuration Properties – C/C++ – Code Generation – Enable Enhanced Instruction Set. Then, include `intrin.h` and `immintrin.h` for at the top. For a list of functions, see here. <https://msdn.microsoft.com/en-us/library/hh977022.aspx>

E Reading Assembly

Reading the assembly code generated by the compiler might help understanding the performance that you get. In GCC, option `-S` gives you the assembly, use `-g -fverbose-asm`