

Performance Modeling: 2D Convolution

Erik Saule

The purpose of the homework is for you to

- build a performance model of a real application
- measure performance of code and compare to the model
- interpret runtime measurement to debug performance issues

1 Generalities

2D Convolution serves as a fundamental operation for a broad range of computer vision tasks. Depending on the values of a kernel, the output of convolution could make pictures to look better by removing noise as well as enabling object recognition by finding edges. Recently, the estimation of kernel values and optimizing machine learning steps have been integrated into one single algorithm called convolutional neural network which has shown a tremendous advancement on image recognition for popular applications such as Google Photos and self-driving car systems.

A convolution \mathbf{C} of dimension k (k is odd) is represented by a 2D array of size $k \times k$ of Single Precision floating point number. Applied on an image \mathbf{in} of size $n \times m$ of Single Precision floating point number, \mathbf{C} creates an image \mathbf{out} of size $n \times m$ of Single Precision floating point number where:

$$out[i][j] = \sum_{x=0}^{k-1} \sum_{y=0}^{k-1} C[x][y] in[i - \left\lfloor \frac{k}{2} \right\rfloor + x][j - \left\lfloor \frac{k}{2} \right\rfloor + y]$$

A page on convolutions: <http://setosa.io/ev/image-kernels/>

Class on convolutional neural network <http://cs231n.stanford.edu/>

(If the band is outside of the image, you can assume the value is the same as the center pixel.)

2 Modeling

Question: How many Flops needs to be done to compute a convolution of dimension k on a image of size $n \times m$?

Question: How much memory needs to be moved to compute a convolution of dimension k on a image of size $n \times m$?

Question: Assuming the performance numbers you measured in assignment 1 and 2, how long should computing a convolution of dimension 3 on an image of 1024×768 take?

Question: What about a convolution of dimension 11 ?

On this kind of problem, performance is usually reported in pixel processed per second.

Question: Plot maximum expected performance as a function of k .

3 Basic Code

Question: Write a basic parallel code to compute convolution of dimension k on an image of size $n \times m$. (Make sure you take n , m , and k as command line parameters.)

Question: Measure and report performance on the following size. Images of 1024×768 , 2048×2048 , 8192×8192 , 4194304×768 and 16777216×768 ; and convolutions of dimension 3, 5, 7, 9, 11, 13, and 15. That is a total of 35 combinations, report measured performances in pixels per second and expected performance.

Question: Why does the performance not always match the model?

4 Optimized Code

Question: Close the performance gap by writing better code and more realistic models.

Question: Measure and report performance of the new code. How close to expected performance can you get?

The typical way to do this is to do iterative loops of model-code-benchmark-analysis.